

---

## GIST: Distributed Training for Large-Scale Graph Convolutional Networks

---

Cameron R. Wolfe<sup>\*1</sup>, Jingkang Yang<sup>\*2</sup>, Arindam Chowdhury<sup>3</sup>, Chen Dun<sup>1</sup>, Artun Bayer<sup>3</sup>,  
Santiago Segarra<sup>3</sup>, and Anastasios Kyrillidis<sup>1</sup>

<sup>1</sup>Department of Computer Science, Rice University, Houston, TX, USA.

<sup>2</sup>School of Computer Science and Engineering, Nanyang Technology University, Singapore.

<sup>3</sup>Department of Electrical and Computer Engineering, Rice University, Houston, TX, USA.

**Main Idea:** We propose graph independent subnetwork training (GIST), a model-agnostic distributed training framework, to pioneer the training of ultra-wide graph convolutional network (GCN) models and reduce wall-clock time of GCN experimentation at arbitrarily large scales. GIST operates by dividing a global GCN model into several, disjoint subnetworks, training them independently, then combining their individual updates back into the global model.

# Background: What is a GCN?

- Deep learning works well for Euclidean data but not all data fits this representation
  - Many applications to rely on graph-structured data (e.g., social networks, chemical molecules, etc.)
- Several deep learning techniques have been extended into the graph domain, the most popular of which is the graph convolutional network (GCN)
  - GCN is a multi-layer architecture that implements a generalization of the convolution operation on graphs
  - GCN works very well for node/graph-level classification tasks

- **The GCN forward pass:**

Aggregate neighboring node representations

Input features: each row is a node feature vector

**Takeaway:** GCNs are just MLPs with intermediate “aggregation” steps between neighboring nodes

$$H_{\ell+1} = \sigma(\bar{A}H_{\ell}\Theta_{\ell}); \text{ where } H_0 = X \text{ and } \bar{A} = D^{-\frac{1}{2}}(A + I)D^{-\frac{1}{2}}$$

Non-linearity  
(e.g., ReLU)

Resembles an MLP by “mixing”  
node representations at each layer

Normalized adjacency matrix  
with added self-loops

# Background: General

- GCNs are the most popular algorithm for deep learning in the graph domain, but their training is notoriously inefficient and scales poorly with graph/model size
  - The “receptive field” of a GCN expands exponentially with the depth (most GCNs have few layers for this reason and because of “over smoothing”)
  - Mini-batches (i.e., subsets of nodes in a graph) are not independent
- *How do we currently tackle this inefficiency?* Partition the graph into multiple sub-graphs that can be used as “mini-batches” during training
  - **Neighborhood sampling:** sample a subset of neighboring nodes at each GCN layer
    - E.g., FastGCN, LADIES, GraphSAGE
  - **Graph partitioning:** divide the graph into smaller sub-graphs for use during training
    - E.g., ClusterGCN, GraphSAINT
- The scale of GCN experimentation still lags behind that of deep learning!
  - Models are small (need more width!) and graph size is limited -- we want to fix this

# Our proposal

- GIST is a novel distributed training methodology that can be used for any GCN architecture and is compatible with neighborhood/graph sampling approaches
- **Methodology:**
  - Randomly decompose global GCN into disjoint, narrow sub-GCNs of equal depth
    - Randomly sample the feature space, each sub-GCNs has the same width
  - Train sub-GCNs independently and in parallel (on separate GPUs)
  - Copy sub-GCN parameters into the global model (disjoint partition prevents collisions)
  - Repeat until convergence
- GIST can be (optionally) combined with node partitioning to handle arbitrarily-large graphs
- **Benefits:** works for any GCN model, improved communication efficiency, improved training efficiency, lower wall-clock time from parallel training, ability to train “ultra-wide” models

# Our proposal

## Algorithm 1 GIST Algorithm

**Parameters:**  $T$  synchronization iterations,  $m$  sub-GCNs,  $\zeta$  local iterations,  $c$  clusters,  $\mathcal{G}$  training graph.

$\Psi_{\mathcal{G}}(\cdot; \Theta) \leftarrow$  randomly initialize GCN

$\{\mathcal{G}_{(j)}\}_{j=1}^c \leftarrow \text{Cluster}(\mathcal{G}, c)$

**for**  $t = 0, \dots, T - 1$  **do**

$\{\Psi_{\mathcal{G}}(\cdot; \Theta^{(i)})\}_{i=1}^m \leftarrow \text{subGCNs}(\Psi_{\mathcal{G}}(\cdot; \Theta), m)$

Distribute each  $\Psi_{\mathcal{G}}(\cdot; \Theta^{(i)})$  to a different worker

**for**  $i = 1, \dots, m$  **do**

**for**  $z = 1, \dots, \zeta$  **do**



$\Psi_{\mathcal{G}}(\cdot; \Theta^{(i)}) \leftarrow \text{subTrain}(\Theta^{(i)}, \{\mathcal{G}_{(j)}\}_{j=1}^c)$

**end for**

**end for**

$\Psi_{\mathcal{G}}(\cdot; \Theta) \leftarrow \text{subAgg}(\{\Theta^{(i)}\}_{i=1}^m)$


**end for**



# Small-Scale Experiments

- **Datasets:** Cora, Pubmed, Citeseer, OGBN-Arxiv
- **Experimental Settings:** 3-layer GCN, width 256, 400 epochs, Adam optimizer
- **Baseline:** identical model trained with standard, single-GPU methodology
- **Findings:** all features except the input can be split, performance is relatively robust to larger numbers of local iterations, performance is relatively robust as the number of sub-GCNs is increased, models trained with GIST often outperform single-GPU baseline models

# Sub-GCNs	$d_0$	$d_1$	$d_2$	Cora	Citeseer	Pubmed	OGBN-Arxiv
Baseline				$81.52 \pm 0.005$	$75.02 \pm 0.018$	$75.90 \pm 0.003$	$70.85 \pm 0.089$
2	✓	✓	✓	$80.00 \pm 0.010$	<b><math>75.95 \pm 0.007</math></b>	$76.68 \pm 0.011$	$65.65 \pm 0.700$
		✓	✓	$78.30 \pm 0.011$	$69.34 \pm 0.018$	$75.78 \pm 0.015$	$65.33 \pm 0.347$
		✓	✓	<b><math>80.82 \pm 0.010</math></b>	$75.82 \pm 0.008$	<b><math>78.02 \pm 0.007</math></b>	<b><math>70.10 \pm 0.224</math></b>
4	✓	✓	✓	$76.78 \pm 0.017$	$70.66 \pm 0.011$	$65.67 \pm 0.044$	$54.21 \pm 1.360$
		✓	✓	$66.56 \pm 0.061$	$68.38 \pm 0.018$	$68.44 \pm 0.014$	$52.64 \pm 1.988$
		✓	✓	<b><math>81.18 \pm 0.007</math></b>	<b><math>76.21 \pm 0.017</math></b>	<b><math>76.99 \pm 0.006</math></b>	<b><math>68.69 \pm 0.579</math></b>
8	✓	✓	✓	$48.32 \pm 0.087$	$45.42 \pm 0.092$	$54.29 \pm 0.029$	$40.26 \pm 1.960$
		✓	✓	$53.60 \pm 0.020$	$54.68 \pm 0.030$	$51.44 \pm 0.002$	$26.84 \pm 7.226$
		✓	✓	<b><math>79.58 \pm 0.006</math></b>	<b><math>75.39 \pm 0.016</math></b>	<b><math>76.99 \pm 0.006</math></b>	<b><math>65.81 \pm 0.378</math></b>



# Large-Scale Experiments

- **Datasets:** Reddit, Amazon2M
- **Experimental Settings:**
  - *Reddit:* GraphSage and GAT architectures, width 256, 2-4 layers, 500 local iterations, partition graph into 15,000 sub-graphs (METIS), batch size 10, Adam optimizer, 80 epochs, 2-8 sub-GCNs
  - *Amazon2M:* GraphSAGE architecture, width 400 or 4096, 5K local iterations, partition graph into 15,000 sub-graphs (METIS), batch size of 10, Adam optimizer, 400 epochs, 2-8 sub-GCNs
- **Baseline:** identical model trained with standard, single-GPU methodology
- **Metrics:** performance is measured as (i) wall-clock time and (ii) F1 classification score

# Large-Scale Experiments

- **Findings on Reddit:** GIST significantly accelerates GCN training and sometimes improves accuracy, more sub-GCNs slightly degrades accuracy but improves wall-clock time, speedup from GIST is more significant for models/datasets with larger compute requirements (e.g., GAT vs. GraphSAGE, Amazon2M vs. Reddit).

<i>L</i>	# Sub-GCNs	GraphSAGE			GAT		
		F1 Score	Time	Speedup	F1 Score	Time	Speedup
2	Baseline	96.09	105.78s	1.00×	89.57	4289.80s	1.00×
	2	96.40	70.29s	1.50×	90.28	2097.16s	2.05×
	4	96.16	68.88s	1.54×	90.02	1112.33s	3.86×
	8	95.46	76.68s	1.38×	89.01	640.44s	6.70×
3	Baseline	96.32	118.37s	1.00×	89.25	7241.23s	1.00×
	2	96.36	80.46s	1.47×	89.63	3432.80s	2.11×
	4	95.76	78.74s	1.50×	88.82	1727.73s	4.19×
	8	94.39	88.54s	1.34×	70.38	944.21s	7.67×
4	Baseline	96.32	120.74s	1.00×	88.36	9969.13s	1.00×
	2	96.01	91.75s	1.32×	87.97	4723.55s	2.11×
	4	95.21	78.74s	1.53×	78.42	2376.21s	4.21×
	8	92.75	88.71s	1.36×	66.30	1262.28s	7.90×



# Large-Scale Experiments

- **Findings on Amazon2M:** narrow models perform worse than the baseline but train much faster, a more significant speedup is observed for wide models (up to >9X), GIST performs better (relative to baseline) with wider models, baseline models take longer to achieve comparable F1 score in comparison to models trained with GIST (e.g., for  $L=2$  and  $m=8$ , standard training takes 2.5X longer to achieve F1 score of 88.86)

$L$	# Sub-GCNs	$d_i = 400$			$d_i = 4096$		
		F1 Score	Time	Speedup	F1 Score	Time	Speedup
2	Baseline	89.90	1.81hr	1.00×	91.25	5.17hr	1.00×
	2	88.36	1.25hr	1.45×	90.70	1.70hr	3.05×
	4	86.33	1.11hr	1.63×	89.49	1.13hr	4.57×
	8	84.73	1.13hr	1.61×	88.86	1.11hr	4.65×
3	Baseline	90.36	2.32hr	1.00×	91.51	9.52hr	1.00×
	2	88.59	1.56hr	1.49×	91.12	2.12hr	4.49×
	4	86.46	1.37hr	1.70×	89.21	1.42hr	6.72×
	8	84.76	1.37hr	1.69×	86.97	1.34hr	7.12×
4	Baseline	90.40	3.00hr	1.00×	91.61	14.20hr	1.00×
	2	88.56	1.79hr	1.68×	91.02	2.77hr	5.13×
	4	87.53	1.58hr	1.90×	89.07	1.65hr	8.58×
	8	85.32	1.56hr	1.93×	87.53	1.55hr	9.13×

# Ultra-Wide GCN Models on Amazon2M

- Ultra-Wide GCN findings:** GIST can train models to SOTA performance with width up to 32K (exceeds capacity of single GPU by ~8X), best performance (91.09 F1 score) is achieved by model with width 16K, standard training methodologies are very limited in model width, more sub-GCNs provide massive speedups as model width increases (e.g., 18X speedup with 2 vs. 8 sub-GCNs with 32K width when L=2), single-GPU models take longer to train cannot match F1 score of ultra-wide models trained with GIST .

<i>L</i> # Sub-GCNs		F1 Score (Time)				
		$d_i = 400$	$d_i = 4096$	$d_i = 8192$	$d_i = 16384$	$d_i = 32768$
2	Baseline	89.38 (1.81hr)	90.58 (5.17hr)	OOM	OOM	OOM
	2	87.48 (1.25hr)	90.09 (1.70hr)	90.87 (2.76hr)	90.94 (9.31hr)	90.91 (32.31hr)
	4	84.82 (1.11hr)	88.79 (1.13hr)	89.76 (1.49hr)	90.10 (2.24hr)	90.17 (5.16hr)
	8	82.56 (1.13hr)	87.16 (1.11hr)	88.31 (1.20hr)	88.89 (1.39hr)	89.46 (1.76hr)
3	Baseline	89.73 (2.32hr)	90.99 (9.52hr)	OOM	OOM	OOM
	2	87.79 (1.56hr)	90.40 (2.12hr)	90.91 (4.87hr)	91.05 (17.7hr)	OOM
	4	85.30 (1.37hr)	88.51 (1.42hr)	89.75 (2.07hr)	90.15 (3.44hr)	OOM
	8	82.84 (1.37hr)	86.12 (1.34hr)	88.38 (1.37hr)	88.67 (1.88hr)	88.66 (2.56hr)
4	Baseline	89.77 (3.00hr)	91.02 (14.20hr)	OOM	OOM	OOM
	2	87.75 (1.79hr)	90.36 (2.77hr)	91.08 (6.92hr)	91.09 (26.44hr)	OOM
	4	85.32 (1.58hr)	88.50 (1.65hr)	89.76 (2.36hr)	90.05 (4.93hr)	OOM
	8	83.45 (1.56hr)	86.60 (1.55hr)	88.13 (1.61hr)	88.44 (2.30hr)	OOM

# Conclusions

- GIST is a distributed training approach for GCNs that enables exploration of larger models and datasets while minimizing the wall-clock time of experiments
  - Compatible with any GCN architecture
  - Can be easily combined with existing sampling approaches for efficient GCN training
- GIST uses a feature-wise partition of a global GCN into several sub-GCNs, which are trained independently and in parallel before having their updates aggregated
- GIST achieves remarkable speed-ups on large-scale GCN experiments and enables the training of ultra-wide GCNs that cannot be handled with standard training methodology
- **GIST is a step towards exploring experimental scales comparable to deep learning within the graph ML community**